

# DSAA3073: Theories in Data Science

## Week 12: Deep Learning Theory: Generalization and Modern Phenomena

---

Concise Learning Sheet

Wei Wang · DSA, HKUST(GZ) · 2026-06-27

Explorative projection

## Opening Challenge

### The Deep Learning Paradox

Modern deep neural networks routinely have billions of parameters—far exceeding the number of training examples. A naive parameter-counting generalization heuristic would suggest an error scale like  $O\left(\sqrt{\frac{p}{n}}\right)$  for  $p$  parameters and  $n$  examples, making the resulting bound **vacuous** (so loose it provides no meaningful information—worse than random guessing). For GPT-3 with 175 billion parameters trained on 300 billion tokens, such a classical-looking estimate tells us essentially nothing.

Yet these networks generalize remarkably well. They achieve state-of-the-art performance on held-out test sets, transfer to new domains, and exhibit emergent capabilities not seen during training.

**The Challenge:** How can we reconcile the success of overparameterized deep learning with classical statistical learning theory? What new principles govern generalization in the modern regime?

**Backbone for Week 12:** This entire week revolves around one central mystery—**why do overparameterized neural networks generalize?** Classical theory says they shouldn't. Every section builds toward resolving this paradox by showing how modern theory extends, refines, and sometimes overturns classical ideas. Keep this question in mind as we progress through each new concept.

# Introduction: The Generalization M

The term “deep learning theory” emerged around 2016-2018 as researchers realized classical tools were insufficient to explain modern networks.

**Motivation:** This week marks the culmination of our journey through statistical learning theory. We’ve built a comprehensive framework spanning **PAC learning** (Probably Approximately Correct learning—a formal framework for when learning is possible, introduced in Weeks 2-3), model selection (Weeks 4-5), kernel methods (Weeks 6-8), boosting and online learning (Weeks 9-10), and deep learning expressivity and optimization (Week 11). Now we confront the central mystery of modern machine learning: **why do overparameterized neural networks generalize?** (where **overparameterized** means having far more parameters than training examples). This question challenges everything we’ve learned about the complexity-generalization tradeoff and forces us to develop new theoretical frameworks that bridge classical statistical learning theory with the empirical success of modern deep learning.

**Prerequisites:** This section reviews concepts from Weeks 2-4. See Week 2 for PAC learning definitions, Week 3 for VC dimension and Rademacher complexity, and Week 4 for bias-variance tradeoff.

## ← From Week 2: PAC Learning

Classical PAC theory tells us that generalization requires hypothesis classes with finite VC dimension, and sample complexity grows with complexity. Yet modern neural networks have VC dimension exceeding the number of training examples—violating the classical regime entirely.

## ← From Week 4: Bias-Variance Tradeoff

The bias-variance decomposition predicts that high-capacity models should have high variance and poor generalization. But overparameterized deep networks seem to escape this tradeoff in the “double descent” regime we’ll explore this week.

## ✓ Checkpoint

**Before we begin:** Given that GPT-3 has 175 billion parameters trained on 300 billion tokens, what does classical VC theory predict about its generalization? Why might this prediction be misleading?

**Answer:** Classical theory would predict poor generalization since parameter count vastly exceeds sample size in naive counting. This prediction is misleading because: (1) effective complexity is lower than parameter count, (2) implicit regularization constrains the solution space, (3) data has structure that reduces intrinsic dimensionality.

## The Classical View vs. Modern Reality

**Motivation:** Before we can understand what makes deep learning special, we need to understand where classical theory succeeds and where it fails. Classical statistical learning theory,

developed primarily in the 1990s-2000s, provides powerful tools for traditional machine learning but breaks down spectacularly for modern deep networks. By examining this breakdown, we'll identify exactly what new theoretical machinery we need.

Classical statistical learning theory rests on three pillars:

1. **Complexity measures: VC dimension** (Vapnik-Chervonenkis dimension—the maximum number of points a hypothesis class can shatter, introduced in Week 2), **Rademacher complexity** (a data-dependent measure of how well a function class can fit random noise, introduced in Week 3)
2. **Bias-variance tradeoff: Underfitting** (high bias—model too simple to capture the true pattern) vs. **overfitting** (high variance—model fits training noise rather than true pattern) (Week 4)
3. **Regularization**: Explicit penalties to control complexity (Week 5)

These tools successfully explain traditional machine learning: SVMs, decision trees, linear models. But they fail spectacularly for deep learning.

#### ← From Week 2: VC Theory

We learned that PAC learnability requires finite VC dimension, and generalization error scales as  $O\left(\sqrt{\frac{d}{n}}\right)$ . This framework beautifully explains when and why simple hypothesis classes generalize.

#### ← From Week 5: Regularization

Classical regularization adds explicit penalties like  $\lambda\|\theta\|^2$  to the loss. But modern networks often use no explicit regularization and still generalize—suggesting a different mechanism is at work.

#### Example: The ImageNet Experiment

Zhang et al. (2017) trained ResNet-50 (25M parameters) on ImageNet (1.3M images):

- **Real labels:** 76% test accuracy (good generalization)
- **Random labels:** 100% training accuracy, 10% test accuracy (pure memorization)

The network has enough capacity to memorize random labels, yet it generalizes on real data. Classical theory cannot distinguish these cases—both have the same VC dimension!

## Roadmap for Week 12

**Motivation:** Week 12 synthesizes the entire course by showing how modern deep learning theory extends, refines, and sometimes overturns classical ideas. We'll see that the story of generalization is richer than classical theory suggests, involving implicit biases, phase transitions, and scaling laws that emerge from the interplay of architecture, optimization, and data structure.

We'll explore four modern phenomena that challenge classical theory:

1. **Classical generalization bounds and their failure** (this section)
2. **Benign overfitting: Interpolation** (achieving zero training error by perfectly fitting all training points) without overfitting (Section 2)
3. **Implicit regularization**: How optimization biases toward simple solutions (Section 3)
4. **Grokking and scaling laws: Delayed generalization** (sudden improvement in test performance long after achieving perfect training accuracy) and **emergent abilities** (capabilities that appear only at sufficient model scale) (Section 4)

Finally, we'll synthesize all 12 weeks into a unified view of learning theory (Section 5).


#### → Week 2 Preview: Double Descent

In Section 2, we'll see how the **double descent** phenomenon resolves the paradox: increasing model capacity beyond the **interpolation threshold** (the point where the model has just enough capacity to perfectly fit training data) can **improve** generalization, contradicting the classical bias-variance curve.

#### → Week 3 Preview: Implicit Regularization

Section 3 reveals that gradient descent has implicit biases—it naturally finds simple, generalizable solutions even without explicit regularization. This explains why overparameterized networks don't overfit despite having no explicit constraints.

### Mental Model: The Four Pillars of Modern Deep Learning Theory



Each classical concept has a modern extension that preserves the core insight while accounting for the unique properties of deep learning.

# Classical Generalization Bounds

**Motivation:** To understand why modern deep learning theory emerged, we must first see where classical bounds fail. Classical generalization bounds—based on VC dimension and Rademacher complexity—provide elegant guarantees for traditional machine learning. But when applied to modern neural networks, these bounds become vacuous, predicting worse-than-random performance for networks that empirically achieve state-of-the-art results. This spectacular failure motivates the new theoretical frameworks we’ll develop in this week.

## ← From Week 2: VC Dimension

In Week 2, we proved that VC dimension characterizes PAC learnability: a class is learnable if and only if it has finite VC dimension. This elegant result works perfectly for linear classifiers, decision trees, and SVMs—but breaks down for deep networks.

## ← From Week 3: Rademacher Complexity

Week 3 introduced Rademacher complexity as a data-dependent refinement of VC dimension. While tighter than VC bounds, Rademacher bounds still fail for modern networks, suggesting the problem is deeper than just loose constant factors.

## VC Dimension Bounds (Week 2 Revisited)

**Motivation:** VC dimension provides the foundation for PAC learning theory—it’s the gold standard for understanding learnability. But when we apply it to neural networks, we discover a shocking result: the bounds are so loose they predict random guessing would outperform state-of-the-art networks. This isn’t just a technical issue—it reveals that VC dimension captures the wrong notion of complexity for deep learning.

Recall the fundamental PAC learning bound:

## ← From Week 2: Fundamental Theorem of PAC Learning

We proved that a hypothesis class is PAC learnable if and only if it has finite VC dimension. This beautiful characterization works perfectly when the VC dimension is much smaller than the sample size ( $d \ll n$ ). But what happens when  $d \gg n$ ?

**Theorem: VC Generalization Bound****Assume:**

- Hypothesis class  $\mathcal{H}$  with  $\text{VCdim}(\mathcal{H}) = d < \infty$
- Loss function  $\ell : \mathcal{H} \times \mathcal{Z} \rightarrow \{0, 1\}$  (binary 0-1 loss)
- Training sample  $S = \{z_1, \dots, z_n\}$  drawn i.i.d. from distribution  $\mathcal{D}$
- Sample size  $n \geq d$  (non-vacuous regime)

**Then:** With probability at least  $1 - \delta$  over the draw of  $S$ , for all  $h \in \mathcal{H}$ :

$$R(h) \leq \hat{R}_S(h) + O\left(\sqrt{\frac{d \log(en/d) + \log(1/\delta)}{n}}\right) \quad (1)$$

where  $R(h) = \mathbb{E}_{z \sim \mathcal{D}}[\ell(h, z)]$  is the true risk and  $\hat{R}_S(h) = \frac{1}{n} \sum_{i=1}^n \ell(h, z_i)$  is the empirical risk.

**Interpretation:** The generalization gap shrinks at rate  $O(\sqrt{d/n})$  when  $d \ll n$ . This bound becomes **vacuous** (larger than 1) when  $d \geq n$ , which is exactly the overparameterized regime where modern deep learning operates.

For neural networks, the VC dimension scales with the number of parameters:

**Proposition: VC Dimension of Neural Networks**

Neural-network VC dimension grows at least linearly with the number of parameters in many standard architectures:

$$\text{VCdim}(\mathcal{H}) = \Omega(p) \quad (2)$$

For many piecewise-polynomial or ReLU network classes, upper bounds are roughly of order  $O(pL \log p)$ , with matching lower bounds depending on architecture and operations. The main lesson is that parameter-count-based VC bounds scale with  $p$  enough to become vacuous for modern networks.

**Pause and Think**

**Question:** What does this bound predict for a network with  $p = 10^9$  parameters and  $n = 10^6$  training examples?

**Hint:** Plug into the VC bound: generalization error  $\approx \sqrt{\frac{10^9}{10^6}} = \sqrt{1000} \approx 30$ .

**Answer:**

The bound predicts generalization error of 30 (or 3000% for classification). This is completely vacuous—random guessing would be better. Yet real networks achieve less than 5% error.

**Rademacher Complexity Bounds (Week 3 Revisited)**

**Motivation:** Rademacher complexity was supposed to be better than VC dimension, often tighter, and accounts for the actual distribution of data. For deep networks, it suffers the same fate: vacuous bounds that tell us little about generalization. Understanding why requires examining how network structure affects complexity.

Bartlett et al. (1998) pioneered norm-based bounds. They're tighter than VC but still often vacuous for deep networks.

Rademacher complexity is data-dependent and often tighter than VC dimension:

### ← From Week 3: Rademacher Complexity

We introduced Rademacher complexity as measuring the ability of a function class to correlate with random noise. For finite classes, this gives tighter bounds than VC dimension. For neural networks, we can bound it using weight norms.

### Definition: Empirical Rademacher Complexity

For function class  $\mathcal{F}$  and sample  $S = \{x_1, \dots, x_n\}$ :

$$\hat{\mathcal{R}}_S(\mathcal{F}) = \mathbb{E}_\sigma \left[ \sup_{f \in \mathcal{F}} \frac{1}{n} \sum_{i=1}^n \sigma_i f(x_i) \right] \quad (3)$$

where  $\sigma_i \in \{-1, +1\}$  are independent **Rademacher random variables** (random signs—each  $\sigma_i$  is  $+1$  or  $-1$  with equal probability).

**Intuition:** This measures how well the function class can correlate with pure random noise. If a class can fit random labels well, it has high complexity and may overfit.

### Theorem: Rademacher Generalization Bound

For bounded losses  $\ell_h(z) \in [0, 1]$ , let  $\mathcal{F} = \{z \rightarrow \ell_h(z) : h \in \mathcal{H}\}$  be the loss-composed class. With probability  $\geq 1 - \delta$ , for all  $h \in \mathcal{H}$ :

$$R(h) \leq \hat{R}_S(h) + 2\hat{\mathcal{R}}_S(\mathcal{F}) + \sqrt{\frac{\log(\frac{1}{\delta})}{2n}} \quad (4)$$

For Lipschitz surrogate losses, contraction inequalities relate the loss-class Rademacher complexity to the prediction-class complexity.

For neural networks, we can bound Rademacher complexity by norms:

### Proposition: Norm-Based Rademacher Bound

For  $L$ -layer network with weight matrices  $W^{(\ell)}$  and **Lipschitz activation**  $\sigma$  (an activation function where  $|\sigma(a) - \sigma(b)| \leq |a - b|$  for all  $a, b$ —examples include ReLU, sigmoid, tanh):

$$\hat{\mathcal{R}}_S(\mathcal{F}) \leq \frac{\prod_{\ell=1}^L \|W^{(\ell)}\|_F}{\sqrt{n}} \quad (5)$$

where  $\|\cdot\|_F$  is the **Frobenius norm** (the square root of the sum of squared entries:  $\|W\|_F = \sqrt{\sum_{i,j} W_{ij}^2}$ ).



**Example: Norm-Based Bound for ResNet**

Consider ResNet-50 with:

- $L = 50$  layers
- Typical weight norms:  $\|W^{(\ell)}\|_F \approx 10$
- Training samples:  $n = 10^6$

The bound gives:

$$\hat{\mathcal{R}}_S \leq \frac{10^{50}}{\sqrt{10^6}} = 10^{47} \quad (6)$$

This is astronomically vacuous! The product of norms explodes exponentially with depth.

## Why Classical Bounds Fail

**Motivation:** Understanding **why** classical bounds fail is more instructive than just observing that they do. Each failure point reveals an implicit assumption that holds for traditional ML but breaks for deep learning. These insights will guide us toward the new theoretical tools we need.

Classical bounds fail for deep learning because they are:

### ← From Week 2: Uniform Convergence

PAC learning theory relies on uniform convergence: if empirical risk approximates true risk uniformly over all hypotheses in  $\mathcal{H}$ , then the empirical risk minimizer generalizes well. This works when  $\mathcal{H}$  is small but fails catastrophically for modern networks.

### ← From Week 10: Online Learning

Week 10 showed that algorithm-specific analysis (regret bounds) can be tighter than worst-case uniform bounds. This same insight applies here: analyzing what gradient descent actually does beats analyzing all possible hypotheses.

### Wrong Path: Worst-Case Analysis

Classical bounds hold for **all** functions in the hypothesis class. They must account for the worst possible function—one that memorizes training data and fails completely on test data.

But gradient descent doesn't find arbitrary functions in  $\mathcal{H}$ . It finds specific functions with implicit structure. The bounds ignore the **optimization algorithm**.

**Wrong Path: Uniform Convergence**

Classical bounds use uniform convergence: bound the gap between empirical and true risk **uniformly** over all  $h \in \mathcal{H}$ .

But we only care about the specific hypothesis  $h_S$  returned by our algorithm on sample  $S$ . Uniform bounds are unnecessarily pessimistic.

**Wrong Path: Ignoring Data Structure**

Classical bounds treat data as arbitrary points in  $\mathcal{X}$ . They don't exploit structure: smoothness, low-dimensional manifolds, natural image statistics.

Real data has rich structure that makes learning easier. The bounds ignore this.

**Wrong Path: Ignoring Implicit Regularization**

Classical bounds assume no regularization beyond explicit penalties. But gradient descent implicitly regularizes: it biases toward simple solutions even without explicit  $\ell_2$  penalties.

This implicit regularization is crucial but invisible to classical theory.

**✓ Checkpoint**

**Self-check:** Which of these four failure modes is most fundamental? Which one, if addressed, would most improve generalization bounds?

**Answer:** Implicit regularization is arguably most fundamental—it explains **why** gradient descent finds simple solutions despite the large hypothesis class. Addressing this leads to algorithmic stability bounds and PAC-Bayes bounds that can be non-vacuous.

**Key Takeaway**

Classical generalization bounds (VC dimension, Rademacher complexity) are **vacuous** for modern deep networks. They fail because they:

1. Analyze worst-case functions, not those found by gradient descent
2. Use uniform convergence over the entire hypothesis class
3. Ignore data structure and implicit regularization

We need new theoretical tools that account for **optimization**, **architecture**, and **data**.

**But this raises deeper questions:** If classical bounds fail so spectacularly, how do we build better theory? We can't simply abandon decades of elegant mathematics. Instead, modern deep learning theory takes four complementary approaches, each addressing a specific failure mode of classical theory. Together, they provide a much richer picture than classical bounds alone—one that actually explains why deep learning works.

## Modern Approaches to Generalization

**Motivation:** If classical bounds fail, we need new tools. Modern deep learning theory doesn't abandon classical ideas—it refines and extends them. Each modern approach addresses one of the failure modes we identified: refined complexity measures address worst-case analysis, PAC-Bayes bounds incorporate the learning algorithm, NTK theory connects to kernel methods where theory is solid, and algorithmic stability accounts for optimization dynamics. Together, these approaches provide a much richer picture than classical theory alone.

How can we do better? Modern deep learning theory takes several complementary approaches:

### → Week 2 Preview: Neural Tangent Kernel

The Neural Tangent Kernel (next section) provides one route to understanding: in the infinite-width limit, neural networks behave like kernel machines, where we have solid theoretical guarantees from Weeks 6-8.

### ← From Week 3: PAC-Bayes

We introduced PAC-Bayes bounds in Week 3 as a way to get data-dependent, algorithm-specific bounds. These become crucial for deep learning, where the choice of initialization and optimization path matters enormously.

### 🧠 Mental Model: Four Routes to Modern Generalization Theory

| Classical Failure                          |   | Modern Solution                               |   | Key Insight                       |
|--------------------------------------------|---|-----------------------------------------------|---|-----------------------------------|
| Parameter counting<br>(VC dimension)       | → | Refined complexity<br>(compression, norms)    | → | Effective<br>dimension<br>matters |
| ↓                                          |   | ↓                                             |   |                                   |
| Worst-case uniform<br>convergence          | → | Algorithm-dependent<br>(PAC-Bayes, stability) | → | What GD<br>finds matters          |
| ↓                                          |   | ↓                                             |   |                                   |
| Arbitrary functions<br>in hypothesis class | → | Kernel regime<br>(NTK theory)                 | → | NTK limit<br>tractable            |
| ↓                                          |   | ↓                                             |   |                                   |
| Explicit<br>regularization only            | → | Implicit<br>regularization                    | → | GD has<br>built-in bias           |

Each approach addresses a specific limitation of classical theory. The real power comes from combining them.

# Benign Overfitting and the Double Descent Phenomenon

**Motivation:** Classical theory tells us that interpolation (perfectly fitting training data) leads to overfitting and poor generalization. This intuition is so ingrained that “overfitting” is synonymous with “fitting too well.” Yet modern deep learning thrives on interpolation—networks with zero training error often generalize best. The “double descent” phenomenon reveals that the relationship between model complexity and generalization is far richer than the classical U-shaped curve suggests. Understanding benign overfitting requires rethinking everything we learned about the bias-variance tradeoff.

Classical wisdom says: models that perfectly fit training data (interpolation) will overfit and generalize poorly. Modern deep learning violates this spectacularly.

**Connection to Week 11:** In Week 11, we examined double descent through the lens of implicit regularization (how gradient descent selects among interpolating solutions). Here we examine it from a generalization theory perspective (how complexity measures change through the interpolation threshold).

## ← From Week 4: Bias-Variance Tradeoff

In Week 4, we learned that test error follows a U-shaped curve: decreasing with complexity (bias reduction) until an optimum, then increasing (variance explosion). Double descent shows this curve has a second descent beyond the interpolation threshold!

## ← From Week 2: Sample Complexity

Classical PAC theory requires  $n \gg d$  (samples exceed VC dimension) for generalization. Benign overfitting occurs when  $p \gg n$  (parameters vastly exceed samples)—the opposite regime!

## ✓ Checkpoint

**Critical question:** If a model achieves 0% training error, what does classical theory predict about its test error? Why might this prediction be wrong?

**Answer:** Classical theory predicts high test error due to overfitting (memorization). This is wrong when: (1) the model class has implicit regularization, (2) there exist simple interpolating solutions, (3) optimization finds those simple solutions.

## The Classical Bias-Variance Tradeoff

Recall from Week 4:

**Theorem: Bias-Variance Decomposition**

For squared loss, at a fixed test point  $x$  and then averaged over  $t$  test error decomposes as:

$$\mathbb{E}[R(\hat{h}_S)] = \text{Bias}^2 + \text{Var} + \sigma^2 \quad (7)$$

where:

- $\text{Bias}^2 = \left(\mathbb{E}[\hat{h}_S(x)] - f^*(x)\right)^2$  (underfitting)
- $\text{Var} = \mathbb{E}\left[\left(\hat{h}_S(x) - \mathbb{E}[\hat{h}_S(x)]\right)^2\right]$  (overfitting)
- $\sigma^2$  is irreducible noise

Belkin et al's 2019 PNAS paper "Reconciling modern machine learning practice and the classical bias-variance tradeoff" sparked intense interest in understanding interpolation.

This predicts a **U-shaped curve**:

**Example: Polynomial Regression**

Fit polynomials of degree  $d$  to  $n = 20$  noisy data points:

- $d = 2$ : High bias, low variance (underfit)
- $d = 5$ : Optimal tradeoff
- $d = 15$ : Low bias, high variance (overfit)
- $d = 20$ : Perfect interpolation, terrible generalization

Classical theory: never interpolate!

**The Double Descent Curve**

Modern experiments reveal a shocking phenomenon:

**Definition: Double Descent Phenomenon**

Test error as a function of model complexity exhibits **two descents**:

1. **Classical regime** ( $p < n$ ): U-shaped curve (bias-variance tradeoff)
2. **Interpolation threshold** ( $p \approx n$ ): Peak in test error (the model has just enough capacity to barely fit training data, often leading to unstable, high-variance solutions)
3. **Modern regime** ( $p \gg n$ ): Decreasing test error despite perfect interpolation!

The curve has a "double descent" shape: 

**Why "double"?** There are two descents: (1) the classical descent from high to low complexity, and (2) the modern descent as we push into massive overparameterization.

**Example: Double Descent in Practice**

Nakkiran et al. (2020) demonstrated double descent across multiple settings:

1. **Model-wise:** Vary network width (ResNet on CIFAR-10)
  - Width 20: Underfit (15% error)
  - Width 50: Optimal (8% error)
  - Width 60: Interpolation threshold (12% error)
  - Width 200: Overparameterized (6% error)
2. **Sample-wise:** Vary training set size (fixed model)
  - Small  $n$ : Overfit
  - $n \approx p$ : Peak error
  - Large  $n$ : Good generalization
3. **Epoch-wise:** Vary training time
  - Early: Underfit
  - Middle: Interpolation threshold
  - Late: Improved generalization (grokking!)

**Benign Overfitting: Theory**

**Motivation:** The term “overfitting” traditionally means “fitting training data too well, leading to poor test performance.” But in the overparameterized regime, we observe **benign overfitting**: the model fits training data perfectly (zero training error) yet still generalizes well. This seems contradictory—how can perfect fitting be benign? The answer lies in **which** perfect-fitting solution the algorithm finds.

How can interpolation generalize? The key is **implicit regularization** toward simple solutions.

**Definition: Minimum-Norm Interpolator**

Among all functions that perfectly fit training data, choose the one with smallest norm:

$$\hat{f} = \operatorname{argmin}_{f \in \mathcal{H}} \|f\|_{\mathcal{H}} \quad \text{subject to} \quad f(x_i) = y_i, \forall i \quad (8)$$

For linear models: this is the minimum  $\ell_2$ -norm solution (the **pseudoinverse solution**  $\hat{\beta} = X^+y$  where  $X^+$  is the Moore-Penrose pseudoinverse). For RKHS: this is the **ridgeless kernel interpolation** solution, the limit of kernel ridge regression as  $\lambda \rightarrow 0$  when the limit is well-defined.

**Key insight:** When there are many ways to achieve zero training error, choosing the simplest one (minimum norm) often generalizes well.

### Author's Take

**Benign overfitting in linear regression (qualitative):** For the minimum-norm interpolator in overparameterized linear regression, interpolation can generalize well under specific covariance spectral-decay/effective-rank conditions, suitable signal assumptions, and controlled noise.

A useful toy intuition is that when  $p \gg n$ , there are many interpolating solutions, and the minimum- $\ell_2$ -norm solution may spread noise across many weak directions rather than fitting it with a large norm. But isotropic features and bounded noise alone are not sufficient for benign overfitting.

### Example: Toy Minimum-Norm Intuition

#### Intuition for $p > n$ :

The pseudoinverse  $\hat{\beta} = X^+y$  (where  $X^+$  is the Moore-Penrose pseudoinverse, which gives the minimum-norm solution when the system is underdetermined) has norm  $\|\hat{\beta}\|^2 = y^T (XX^T)^{-1} y$ .

In favorable random-design models,  $XX^T$  may be well-conditioned and the solution lies in the row space of  $X$ , which is at most  $n$ -dimensional (even though the parameter space is  $p$ -dimensional with  $p \gg n$ ).

Under the right spectral conditions, this effective geometry can keep prediction error near the irreducible noise level even while training error is zero. The precise Bartlett et al. theorem is more delicate than this cartoon and depends on covariance eigenvalue decay and effective ranks.

**Concrete intuition:** Imagine fitting  $n = 10$  noisy points with  $p = 100$  parameters. There are infinitely many solutions with zero training error. The minimum-norm solution spreads the weights across all parameters gently, rather than using a few large weights. This spreading prevents overfitting to noise.

### Pause and Think

**Question:** Why does increasing  $p$  beyond  $n$  **improve** generalization?

**Hint:** More parameters give more freedom to find a simple interpolator. The minimum-norm constraint becomes easier to satisfy.

**Answer:**

With many parameters, there are many ways to interpolate. Under favorable spectral structure, minimum-norm interpolation can choose a low-complexity solution that generalizes well. This is implicit regularization at work!

## Conditions for Benign Overfitting

**Motivation:** Benign overfitting is not automatic—it requires specific conditions on the data, model, and optimization. Understanding when it works and when it fails helps us design better learning systems.

Benign overfitting doesn't always occur. It requires:

1. **Overparameterization:**  $p \gg n$  (not just  $p > n$ )
2. **Implicit regularization:** Optimization biases toward simple solutions
3. **Data structure:** Features have appropriate spectral properties
4. **Low noise:** Signal-to-noise ratio is favorable

#### Wrong Path: Benign Overfitting is Universal

Benign overfitting is **not** universal. It can fail when:

- Data is adversarial or has heavy-tailed noise
- Features are poorly conditioned
- Optimization doesn't find minimum-norm solution

Example: Random Fourier features with uniform spectrum show **harmful** overfitting even when  $p \gg n$ .

#### Connection

**Connection to Week 6:** The minimum-norm interpolator in an RKHS is the ridgeless limit of kernel ridge regression under appropriate conditions. Benign overfitting helps explain why kernel methods can sometimes interpolate and still generalize!

**Connection to Week 4:** The bias-variance tradeoff still holds, but in the overparameterized regime, variance **decreases** with more parameters (counterintuitively). This is because the minimum-norm constraint becomes more effective.

#### Key Takeaway

**Benign overfitting** shows that perfect interpolation can generalize well in the overparameterized regime ( $p \gg n$ ). The key mechanisms are:

1. Minimum-norm implicit regularization
2. Effective dimension is governed by covariance spectrum and sample geometry, not raw parameter count alone
3. Overparameterization provides flexibility to find simple interpolators

This reconciles modern deep learning practice with statistical learning theory.

**Progress on backbone:** Double descent reveals that overparameterization can help generalization—but it doesn't explain **why**. Among infinitely many solutions that perfectly fit training data, why does gradient descent consistently find the ones that generalize? We've mentioned "implicit regularization" several times, but now we're ready to make it precise. This section reveals the hidden hand that guides optimization toward simple, generalizable solutions.



# Implicit Regularization: The Hidden

This theorem shows gradient descent implicitly maximizes the margin—exactly what SVMs do explicitly! The connection to Week 9 (boosting and margins) is profound.

**Motivation:** One of the deepest mysteries in deep learning is this: why do neural networks generalize when trained with gradient descent on unregularized loss functions? There's no explicit  $\lambda\|\theta\|^2$  penalty, no dropout during training in many cases, yet the networks don't overfit. The answer: gradient descent itself acts as a regularizer, implicitly biasing toward certain solutions over others. This **implicit regularization** is invisible in the loss function but emerges from the optimization dynamics.

Why does gradient descent find minimum-norm solutions? There's no explicit  $\|\theta\|^2$  penalty in the loss function. The answer: **implicit regularization**.

## Implicit Bias in Linear Models

**Motivation:** To build intuition for implicit regularization, we start with the simplest case: linear models trained on separable data. Here, we can prove exactly what gradient descent does: it implicitly maximizes the margin, recovering the SVM solution without any explicit max-margin objective.

Let's start with the simplest case: linear models with separable data.

### Theorem: Implicit Bias of Gradient Descent (Soudry et al. 2018)

#### Assume:

- Linear model  $f(x, \theta) = \langle \theta, x \rangle$  for binary classification
- Training data  $\{(x_i, y_i)\}_{i=1}^n$  is linearly separable:  $\exists \theta^*$  with  $y_i \langle \theta^*, x_i \rangle > 0$  for all  $i$
- Loss function: exponentially-tailed (e.g., logistic loss  $\ell(\text{margin}) = \log(1 + e^{-\text{margin}})$ )
- Gradient descent (full-batch, not SGD) with sufficiently small constant learning rate  $\eta$
- Initialization near zero:  $\|\theta(0)\| = o(1)$

**Then:** As training progresses with  $t \rightarrow \infty$ :

1. Loss converges:  $\mathcal{L}(\theta(t)) \rightarrow 0$
2. Norm diverges:  $\|\theta(t)\| \rightarrow \infty$  logarithmically (grows like  $\log t$ )
3. Direction converges to max-margin separator:

$$\lim_{t \rightarrow \infty} \frac{\theta(t)}{\|\theta(t)\|} = \frac{\theta_{\text{SVM}}}{\|\theta_{\text{SVM}}\|} \quad (9)$$

where  $\theta_{\text{SVM}} = \arg \max_{\theta: y_i \langle \theta, x_i \rangle \geq 1 \forall i} \frac{1}{\|\theta\|}$  is the hard-margin SVM solution.

**Interpretation:** Gradient descent implicitly regularizes toward large-margin solutions even without explicit  $\ell_2$  penalty. This explains why overparameterized linear models don't overfit despite having zero explicit regularization.

#### Proof. Sketch:

Consider gradient flow  $\frac{d\theta}{dt} = -\nabla \mathcal{L}(\theta)$  with exponentially-tailed loss (e.g., logistic  $\ell(\text{margin}) = \log(1 + e^{-\text{margin}})$ ).

**Step 1 (Weight Divergence):** For linearly separable data, zero loss requires infinite margins. Since  $\mathcal{L}(\theta) \rightarrow 0$  implies  $y_i \langle \theta, x_i \rangle \rightarrow \infty$  for all  $i$ , we have  $\|\theta\| \rightarrow \infty$ . The growth rate is logarithmic:  $\|\theta(t)\| = \Theta(\log t)$ .

**Step 2 (Direction Convergence):** Decompose  $\theta(t) = r(t)\hat{\theta}(t)$  where  $r(t) = \|\theta(t)\|$  and  $\|\hat{\theta}(t)\| = 1$ . The gradient is dominated by examples with smallest margins:

$$\nabla \mathcal{L}(\theta) \approx - \sum_{i: \text{small margin}} y_i x_i e^{-y_i \langle \theta, x_i \rangle} \quad (10)$$

As  $r \rightarrow \infty$ , the direction dynamics satisfy:

$$\frac{d\hat{\theta}}{dt} \approx -\nabla \left( \min_i y_i \langle \hat{\theta}, x_i \rangle \right) \quad (11)$$

This converges to  $\arg \max_{\hat{\theta}} \min_i y_i \langle \hat{\theta}, x_i \rangle$  (the max-margin direction).

**Step 3 (SVM Connection):** The max-margin direction is exactly the hard-margin SVM solution:

$$\hat{\theta}_{\text{SVM}} = \arg \max_{\hat{\theta}: y_i \langle \hat{\theta}, x_i \rangle \geq 1 \forall i} \frac{1}{\|\hat{\theta}\|} \quad (12)$$

**Extension to Deep Networks:** In the NTK regime (wide networks, small learning rates), the network dynamics linearize around initialization, and the same implicit bias argument applies in the feature space defined by the Neural Tangent Kernel (see Week 11).

For rigorous proofs, see Soudry et al. (2018) and Lyu & Li (2020). ■

#### Example: Visualizing Implicit Bias

Train logistic regression on 2D separable data:

- **Epoch 10:** Loss = 0.5,  $\|\theta\| = 2$ , margin = 0.3
- **Epoch 100:** Loss = 0.01,  $\|\theta\| = 5$ , margin = 0.8
- **Epoch 1000:** Loss =  $10^{-6}$ ,  $\|\theta\| = 8$ , margin = 0.95
- **Epoch 10000:** Loss =  $10^{-12}$ ,  $\|\theta\| = 11$ , margin = 0.99

The direction  $\frac{\theta}{\|\theta\|}$  converges to the max-margin separator!

#### Connection

**Connection to Week 9:** AdaBoost also maximizes the margin! Both gradient descent and boosting exhibit the same implicit bias toward large-margin solutions.

**Unifying principle:** Iterative optimization with exponentially-tailed losses implicitly maximizes margins, explaining generalization without explicit regularization.

## Implicit Bias in Deep Networks

**Motivation:** For deep networks, implicit regularization is more subtle than the linear case. We can't simply say "it converges to the max-margin solution." Instead, recent work has identified

several mechanisms: training at the edge of stability, preferring flat biases. These mechanisms interact in complex ways, but all contribute to overfitting.

Cohen et al. (2021) discovered the edge of stability phenomenon. It's a striking example of emergent behavior in gradient descent.

For deep networks, implicit regularization is more complex but equally important.

**Connection to Week 11:** We introduced the edge-of-stability phenomenon in Week 11 (gradient descent with learning rates near the stability threshold  $2/L$ ). Here we examine its implications for generalization theory and implicit regularization.

#### Definition: Edge of Stability Phenomenon

The **edge-of-stability** phenomenon refers to empirical observations where, during training, the maximum Hessian eigenvalue  $\lambda_{\max}(\nabla^2 \mathcal{L})$  (the largest eigenvalue of the loss Hessian—a measure of curvature) often hovers near the gradient-descent stability threshold:

$$\lambda_{\max} \approx \frac{2}{\eta} \quad (13)$$

where  $\eta$  is the learning rate. This behavior appears in some settings, but it is not a universal theorem.

**Intuition:** For gradient descent to be stable, we need  $\eta \lambda_{\max} < 2$ . The “edge of stability” observation is that training often self-organizes to keep  $\lambda_{\max} \approx \frac{2}{\eta}$ —right at the boundary between stable and unstable dynamics.

#### Example: Edge of Stability in Practice

Train ResNet-18 on CIFAR-10 with learning rate  $\eta = 0.1$ :

- **Epoch 1-10:**  $\lambda_{\max}$  grows from 5 to 20
- **Epoch 10-50:**  $\lambda_{\max}$  oscillates around  $\frac{2}{\eta} = 20$
- **Epoch 50-100:**  $\lambda_{\max}$  remains near 20 despite changing loss landscape

The system may appear to maintain  $\lambda_{\max} \approx \frac{2}{\eta}$  through a self-stabilization mechanism in this setting.

#### Author's Take

**Sharpness and implicit regularization:** Appropriately normalized notions of flatness often correlate with generalization, and large learning rates can bias optimization toward wider basins in some regimes.

Naive Hessian sharpness is not parameterization invariant, and the relation  $\lambda_{\max} \approx \frac{2}{\eta}$  is best viewed as an empirical edge-of-stability pattern rather than a universal theorem.

This explains why:

- Larger learning rates may bias toward flatter minima in some regimes, but too-large rates can destabilize training
- Learning rate warmup helps: start small (explore), then increase (flatten)
- Learning rate decay at the end: settle into a flat minimum

**Pause and Think**

**Question:** Why do flat minima generalize better than sharp minima?

**Hint:** Flat minima are robust to perturbations. Sharp minima are sensitive to small changes in parameters or data.

**Answer:**

A flat minimum represents a “wide valley” in the loss landscape. Small perturbations (from noise, data variation, or parameter changes) don’t significantly affect the loss. This robustness translates to better generalization.

## Sharpness-Aware Minimization (SAM)

**Motivation:** If flat minima generalize better (as empirical evidence suggests), can we explicitly search for them? Sharpness-Aware Minimization (SAM) does exactly this by modifying gradient descent to seek parameters where the loss remains low even under small perturbations.

Can we explicitly seek flat minima? Yes!

**Definition: Sharpness-Aware Minimization**

Instead of minimizing loss  $\mathcal{L}(\theta)$ , minimize the **worst-case loss in a neighborhood**:

$$\min_{\theta} \max_{\|\varepsilon\| \leq \rho} \mathcal{L}(\theta + \varepsilon) \quad (14)$$

This explicitly seeks flat minima where the loss doesn’t increase much under perturbations.

**Intuition:** A sharp minimum is one where small changes to  $\theta$  cause large changes in loss. By minimizing the worst-case loss in a ball of radius  $\rho$ , we force the algorithm to find regions where the loss is flat—the loss stays low even under perturbation.

**Example: SAM Algorithm**

At each step:

1. Compute adversarial perturbation:  $\varepsilon = \rho \frac{\nabla \mathcal{L}(\theta)}{\|\nabla \mathcal{L}(\theta)\|}$
2. Compute gradient at perturbed point:  $g = \nabla \mathcal{L}(\theta + \varepsilon)$
3. Update:  $\theta \leftarrow \theta - \eta g$

SAM consistently improves generalization across architectures and datasets!

### Connection

**Connection to Week 5:** SAM is related to adversarial training and robust optimization. It's an explicit form of the implicit regularization that gradient descent naturally performs.

**Connection to PAC-Bayes:** Flat minima have small PAC-Bayes bounds because the posterior  $Q$  can be wide (high entropy) while staying close to the prior  $P$ .

## Implicit Regularization Mechanisms

Multiple mechanisms contribute to implicit regularization:

### Definition: Sources of Implicit Regularization

1. **Algorithmic:**
  - SGD noise (mini-batch sampling)
  - Learning rate schedule
  - Optimization algorithm (Adam, momentum, etc.)
2. **Architectural:**
  - Depth (compositional structure)
  - Skip connections (ResNets)
  - Normalization (BatchNorm, LayerNorm)
  - Dropout
3. **Initialization:**
  - Scale (Xavier, He, NTK)
  - Distribution (Gaussian, orthogonal)
4. **Data:**
  - Augmentation (flips, crops, mixup)
  - Structure (manifold hypothesis)

### Example: BatchNorm as Implicit Regularization

Batch normalization:

$$\hat{x} = \frac{x - \mu_B}{\sqrt{\sigma_B^2 + \varepsilon}} \quad (15)$$

implicitly regularizes by:

1. Reducing sensitivity to weight scale (scale invariance)
2. Adding noise through batch statistics
3. Smoothing the loss landscape

Networks with BatchNorm generalize better even without explicit regularization!

### Key Takeaway

**Implicit regularization** is the hidden mechanism that makes deep learning work:

1. Gradient descent biases toward max-margin solutions (linear case)
2. Training operates at the edge of stability (deep networks)
3. Flat minima generalize better than sharp minima
4. Multiple sources: algorithm, architecture, initialization, data

This explains why overparameterized networks don't overfit despite having no explicit regularization.

### Confidence Check

**Self-Assessment: Can you...**

1. Explain why gradient descent on separable data converges to the max-margin solution?
2. Describe at least three sources of implicit regularization in modern networks?
3. Connect implicit regularization to PAC-Bayes bounds (how does staying near initialization help)?
4. Explain the edge-of-stability phenomenon and why it relates to generalization?

**If you can answer all four:** You understand implicit regularization deeply. This is one of the most important concepts in modern deep learning theory!

**If you struggle with any:** Review Section 3, especially the implicit bias theorem and the connection to margin maximization. The key insight is that optimization itself acts as regularization.

**Having understood the mechanisms, we can finally complete the picture:** We've explored why classical bounds fail (Section 1), how modern approaches address these failures (Section 2), why interpolation can be benign (Section 3), and how implicit regularization guides optimization (Section 4). Now we're ready to see the payoff: modern generalization bounds that are actually non-vacuous for real neural networks. These bounds synthesize everything we've learned into practical guarantees.

## Modern Generalization Bounds

**Motivation:** Classical bounds failed because they were too pessimistic—analyzing worst-case functions in huge hypothesis classes. Modern bounds take a different approach: they account for the learning algorithm (PAC-Bayes), the effective complexity of learned functions (compression), the geometric properties of solutions (norm-based bounds), and the stability of the optimization process. These bounds can be non-vacuous for real networks, finally bridging theory and practice.

Armed with insights about implicit regularization, we can derive tighter generalization bounds.

### ← From Week 3: PAC-Bayes Framework

PAC-Bayes bounds from Week 3 now become our primary tool. By choosing the prior  $P$  at initialization and posterior  $Q$  after training, we can measure how much the algorithm has moved—capturing implicit regularization mathematically.

### Connection

**Looking ahead:** The open problems section will show that even these modern bounds are not tight enough. Closing the theory-practice gap remains one of the grand challenges of machine learning theory.

## PAC-Bayes Bounds for Neural Networks

### Theorem: PAC-Bayes Bound (Refined)

#### Assume:

- Hypothesis space  $\mathcal{H}$  with prior distribution  $P$  over  $\mathcal{H}$  (fixed before observing data)
- Loss function  $\ell : \mathcal{H} \times \mathcal{Z} \rightarrow [0, 1]$  (bounded)
- Training sample  $S = \{z_1, \dots, z_n\}$  drawn i.i.d. from distribution  $\mathcal{D}$
- Posterior distribution  $Q$  over  $\mathcal{H}$  (can depend on  $S$ )

**Then:** With probability at least  $1 - \delta$  over the draw of  $S$ , for all posteriors  $Q$ :

$$R(Q) \leq \hat{R}_S(Q) + \sqrt{\frac{\text{KL}(Q \parallel P) + \log(2\sqrt{n}/\delta)}{2(n-1)}} \quad (16)$$

where  $R(Q) = \mathbb{E}_{\theta \sim Q}[R(\theta)]$  is the expected true risk and  $\hat{R}_S(Q) = \mathbb{E}_{\theta \sim Q}[\hat{R}_S(\theta)]$  is the expected empirical risk under the posterior distribution  $Q$ .

**Interpretation:** The generalization gap depends on how much the posterior  $Q$  differs from the prior  $P$  (measured by KL divergence). Small deviation from prior beliefs leads to better generalization. This bound naturally handles stochastic predictors and captures implicit regularization through the KL term.

**Example: Computing PAC-Bayes Bounds**

Dziugaite & Roy (2017) computed non-vacuous bounds for MNIS

1. **Prior  $P$ :** Gaussian at initialization  $\mathcal{N}(\theta_0, \sigma_0^2 I)$
2. **Posterior  $Q$ :** Gaussian around trained weights  $\mathcal{N}(\theta^*, \sigma^2 I)$
3. **KL divergence:** for isotropic Gaussians, use the standard Gaussian KL formula including mean, trace, log-determinant, and dimension terms; in the equal-variance case this reduces mainly to a squared-distance term.

Result: Test error  $\leq 5.8\%$  (bound) vs.  $2.1\%$  (actual)—the first non-vacuous bound for neural networks!

The key insight: if training doesn't move far from initialization (small  $\|\theta^* - \theta_0\|$ ), the KL term is small and the bound is tight.

**Pause and Think**

**Question:** How can we make PAC-Bayes bounds tighter?

**Hint:** Reduce  $\text{KL}(Q \parallel P)$  by: (1) staying close to initialization, (2) using a better prior, or (3) using a more concentrated posterior.

**Answer:**

Practical strategies: (1) early stopping, (2) pre-training as prior, (3) pruning to reduce effective dimensionality. All exploit implicit regularization!

**Compression-Based Bounds**

If a network can be compressed without losing accuracy, its effective complexity is small.

**Theorem: Compression Bound**

**Assume:**

- Training sample  $S = \{z_1, \dots, z_n\}$  drawn i.i.d. from distribution  $\mathcal{D}$
- Loss function  $\ell : \mathcal{H} \times \mathcal{Z} \rightarrow [0, 1]$  (bounded)
- Hypothesis  $h$  can be represented using  $k$  bits (after quantization, pruning, or other compression)
- Compression scheme is data-independent (does not peek at test data)

**Then:** With probability at least  $1 - \delta$  over the draw of  $S$ :

$$R(h) \leq \hat{R}_S(h) + O\left(\sqrt{\frac{k + \log(1/\delta)}{n}}\right) \quad (17)$$

**Interpretation:** The bound depends on  $k$  (compressed size in bits), not  $p$  (original parameter count). If a network with billions of parameters can be compressed to a few megabytes without losing training accuracy, the effective complexity is determined by the compressed size, explaining why overparameterization doesn't hurt generalization.



**Example: Compression in Practice**

Han et al. (2015) compressed AlexNet:

- Original: 240 MB (61M parameters  $\times$  32 bits)
- Pruned: 90% weights removed gives about 6.1M remaining parameters
- Quantized: 8 bits per weight
- Compressed: 6.9 MB ( $35\times$  compression)
- Accuracy loss: less than 1%

Effective complexity:  $k = 6.1M \times 8 = 49M$  bits vs. original  $p = 61M \times 32 = 2B$  bits— $40\times$  reduction!

**Connection**

**Connection to Week 2:** Compression bounds are related to Occam's razor and minimum description length (MDL). Simpler hypotheses (shorter descriptions) generalize better.

**Connection to Week 5:** Pruning is a form of  $\ell_0$  regularization. Compression bounds formalize why sparse models generalize well.

**Norm-Based Bounds (Refined)**

Classical norm-based bounds multiply layer norms, leading to exponential blow-up. Modern bounds are more sophisticated.

**Remark****Spectral-Norm Generalization Bounds (Bartlett et al. 2017, Neyshabur et al. 2018)**

For deep networks, Rademacher complexity can be bounded using spectral norms:

$$\mathcal{R}_n(\mathcal{F}) \leq \frac{C}{n} \cdot L \cdot \left( \prod_{\ell=1}^L \|W^{(\ell)}\|_2 \right) \cdot (\text{margin factors}) \quad (18)$$

where  $\|W^{(\ell)}\|_2$  is the spectral norm (largest singular value) of layer  $\ell$ ,  $L$  is the number of layers, and margin factors depend on the minimum margin achieved on the training data.

**Why this helps:** Spectral norms  $\|W\|_2$  are often much smaller than Frobenius norms  $\|W\|_F$  when weight matrices have low effective rank. For overparameterized networks,  $\|W\|_2 \ll \sqrt{p}$  even when  $\|W\|_F \approx \sqrt{p}$ , leading to non-vacuous bounds in some settings.

**Key improvement over classical bounds:** Spectral information and margins can be much tighter than raw parameter count or products of Frobenius norms. However, these bounds still involve products over all layers and can be loose for very deep networks.

For precise statements with all constants and conditions, see the original papers.

**Example: Spectral Normalization**

Miyato et al. (2018) proposed **spectral normalization**: constraining the spectral norm of the weight matrices during training.

Benefits:

1. Tighter generalization bounds
2. Lipschitz continuity (robustness)
3. Stable training (especially for GANs)

With additional norm-ratio and margin control, the bound can become much less dependent on width, though precise constants and factors matter.

Spectral normalization is widely used in GANs (Wasserstein GANs, StyleGAN) to ensure stable training and good generalization.

## Algorithmic Stability Bounds

Stability theory provides algorithm-dependent bounds.

**Theorem: Stability of SGD (Hardt et al. 2016)**

For bounded,  $L$ -Lipschitz,  $\beta$ -smooth losses under the standard stability assumptions, SGD with learning rate  $\eta \leq \frac{1}{\beta}$  and  $T$  steps has stability scaling like:

$$\varepsilon = O\left(\frac{L^2 \eta T}{n}\right) \quad (19)$$

This gives a corresponding expected generalization bound, with high-probability versions requiring additional concentration terms.

**Example: Stability in Practice**

For a toy stability calculation for ResNet-50 on ImageNet, suppose:

- Lipschitz constant:  $L \approx 10$
- Learning rate:  $\eta = 0.1$
- Training steps:  $T = 10^5$
- Sample size:  $n = 10^6$

Stability bound:  $\varepsilon \approx \frac{2 \times 100 \times 0.1 \times 10^5}{10^6} = 0.2$

This toy number is non-vacuous, but certified worst-case Lipschitz constants for real networks can be much larger, so practical stability bounds must be interpreted carefully.

### Connection

**Connection to Week 3:** We introduced stability for leave-one-out cross-validation. Now we see it's fundamental for understanding SGD's generalization.

**Connection to Week 10:** Online learning algorithms have regret bounds that imply stability. The connection: online learning to low regret to stability to generalization.

## Comparison of Modern Bounds

| Bound Type  | Key Quantity                     | Pros                                      | Cons                              |
|-------------|----------------------------------|-------------------------------------------|-----------------------------------|
| PAC-Bayes   | $KL(Q \parallel P)$              | Data-dependent, accounts for optimization | Requires choosing prior/posterior |
| Compression | Compressed size $k$              | Intuitive, practical                      | Requires compression algorithm    |
| Norm-based  | Product of norms                 | Architecture-dependent                    | Can be loose for deep networks    |
| Stability   | Stability constant $\varepsilon$ | Algorithm-dependent                       | Requires smoothness assumptions   |

Table 1: Comparison of modern generalization bounds

### Key Takeaway

Modern generalization bounds improve on classical VC/Rademacher bounds by:

1. **PAC-Bayes:** Account for the learning algorithm via KL divergence
2. **Compression:** Measure effective complexity, not raw parameter count
3. **Norm-based:** Use spectral norms and margins for tighter bounds
4. **Stability:** Directly analyze the optimization algorithm

These bounds can be non-vacuous in selected practical settings, unlike many classical bounds. However, many remain loose for large modern networks and still don't fully explain deep learning's success—open problems remain!

## Summary: Modern Generalization Theory

We've seen four complementary approaches to understanding generalization in deep learning:

1. **Neural Tangent Kernel:** Connects infinite-width networks to kernel methods, explains lazy training regime but misses feature learning
2. **Benign Overfitting:** Shows interpolation can generalize in the overparameterized regime via minimum-norm implicit regularization
3. **Implicit Regularization:** Gradient descent biases toward simple solutions (max-margin, flat minima) without explicit penalties
4. **Modern Bounds:** PAC-Bayes, compression, norm-based, and stability bounds that can be non-vacuous in selected settings

### Connection

**Synthesis:** These approaches are complementary, not competing:

- NTK explains **what** networks learn in the infinite-width limit
- Benign overfitting explains **why** interpolation works
- Implicit regularization explains **how** optimization finds good solutions
- Modern bounds **quantify** generalization guarantees

Together, they provide a much richer understanding than classical theory alone.

**Final backbone resolution:** We began Week 12 asking: “Why do overparameterized neural networks generalize?” We can now answer with precision:

1. **Classical theory fails** because it ignores optimization, architecture, and data structure (Section 1)
2. **Modern approaches** refine complexity measures and account for the learning algorithm (Section 2)
3. **NTK theory** shows infinite-width networks behave like kernel machines with rigorous guarantees (Section 3)
4. **Benign overfitting** reveals that interpolation doesn't cause overfitting when implicit regularization is strong (Section 4)
5. **Implicit regularization** explains why gradient descent finds simple solutions without explicit penalties (Section 5)
6. **Modern bounds** quantify these insights into non-vacuous guarantees (Section 6)

The story is complete: overparameterization provides flexibility, implicit regularization constrains the search, and the combination enables both perfect training fit and excellent generalization.

### Thought Experiment

**Thought experiment:** Imagine training two networks on the same data:

1. Network A: Standard initialization, SGD with learning rate 0.1
2. Network B: Same architecture, but initialized at a trained network from a different task

Which generalizes better? Why?

**Answer:** Network B likely generalizes better because:

- The prior  $P$  (pre-trained network) is better than random initialization
- $\text{KL}(Q \parallel P)$  is smaller (less movement from prior)
- PAC-Bayes bound is tighter


This explains why transfer learning and pre-training work so well!

## Course Synthesis: A Unified View of Learning Theory

**Motivation:** This final section is where everything comes together. Over 12 weeks, we’ve built a comprehensive theory of learning—from the foundational question “what is learnable?” to the modern mystery of deep learning generalization. Each week added a piece to the puzzle. Now we step back to see the complete picture: how PAC learning, kernel methods, boosting, online learning, and deep learning theory form a coherent whole. This synthesis reveals that seemingly disparate results are manifestations of a few fundamental principles about the nature of learning from data.

We’ve reached the end of our 12-week journey through statistical learning theory. Let’s synthesize everything into a unified framework.

### Mental Model: The Learning Theory Landscape



The course forms a cycle: we start with fundamental questions about learnability, develop sophisticated tools (kernels, boosting, online learning), and return to fundamental questions with new insights (deep learning). Each stage deepens our understanding.

## The Arc of DSAA3073

### Connection

#### The Five Themes:

1. **Foundations (Weeks 1-3):** PAC learning, VC dimension, Rademacher complexity
2. **Classical Theory (Weeks 4-5):** Bias-variance, regularization, model selection
3. **Kernel Methods (Weeks 6-8):** RKHS, kernel trick, representer theorem
4. **Ensemble & Online (Weeks 9-10):** Boosting, margin theory, regret bounds
5. **Modern Deep Learning (Weeks 11-12):** Expressivity, NTK, implicit regularization

Each theme builds on the previous, culminating in a comprehensive understanding of learning.

## Unifying Principles

Despite the diversity of topics, several principles unify the entire course:

### Principle 1: The Generalization-Complexity Tradeoff

**Motivation:** The tension between fitting training data and generalizing to new data is the central theme of learning theory. From VC dimension to PAC-Bayes KL divergence, every week introduced a different way to measure complexity. But they all tell the same story: generalization requires controlling complexity, and the optimal complexity depends on the amount of data available.

From Week 2 to Week 12, we've seen different manifestations of the same fundamental tradeoff:

#### ← From Week 2: VC Dimension

The VC dimension was our first complexity measure. It counts the maximum number of points that can be shattered (all labelings realized). Simple insight: larger classes fit more patterns but need more data to generalize.

#### ← From Week 4: Bias-Variance

The bias-variance decomposition quantifies the tradeoff numerically: complex models have low bias but high variance. The optimal complexity balances these competing forces.

#### ← From Week 6: RKHS Norm

In reproducing kernel Hilbert spaces, complexity is measured by  $\|f\|_{\mathcal{H}}$ . The representer theorem shows that minimizing training error plus RKHS norm yields sparse solutions—automatic complexity control!

| Week | Complexity Measure                        | Generalization Bound                                                        |
|------|-------------------------------------------|-----------------------------------------------------------------------------|
| 2-3  | VC dimension $d$                          | $O\left(\sqrt{\frac{d}{n}}\right)$                                          |
| 3    | Rademacher complexity $\hat{\mathcal{R}}$ | $O\left(\hat{\mathcal{R}} + \sqrt{\frac{\log(\frac{1}{\delta})}{n}}\right)$ |
| 4    | Model complexity                          | Bias-variance tradeoff                                                      |
| 6-8  | RKHS norm $\ f\ _{\mathcal{H}}$           | Representer theorem                                                         |
| 9    | Margin $\gamma$                           | $O\left(\frac{1}{\gamma\sqrt{n}}\right)$                                    |
| 12   | PAC-Bayes KL                              | $O\left(\sqrt{\text{KL} \frac{Q\ P}{n}}\right)$                             |

Table 2: Evolution of complexity measures across the course

**Unifying insight:** Generalization requires controlling complexity. The form of complexity depends on the setting, but the principle is universal.

## Principle 2: From Worst-Case to Average-Case

**Motivation:** One of the most important progressions in the course is from worst-case to average-case analysis. Worst-case bounds (like VC theory) must account for pathological data distributions and adversarial examples. Average-case bounds exploit the structure in real data—smoothness, manifolds, natural statistics. This progression reflects the maturation of learning theory from pure mathematics to practical science.

The course progresses from worst-case to average-case analysis:

### ← From Week 2: VC Dimension

VC dimension provides distribution-free guarantees—the bound holds for **any** distribution. This generality comes at a cost: the bounds are loose because they can't exploit structure in the specific distribution we care about.

### ← From Week 3: Rademacher Complexity

Rademacher complexity is our first step toward average-case analysis. It depends on the actual sample  $S$  drawn from the distribution, not just the hypothesis class  $\mathcal{H}$ . This allows tighter bounds that exploit data geometry.

### ← From Week 10: Regret Bounds

Online learning completely abandons distributional assumptions. Regret bounds hold even for adversarially chosen sequences! Yet by analyzing the actual sequence (not worst-case over all sequences), we get algorithm-specific guarantees.

- **Week 2:** VC dimension (worst-case over all functions in  $\mathcal{H}$ )
- **Week 3:** Rademacher complexity (average over data distribution)
- **Week 10:** Regret bounds (average over sequence of examples)
- **Week 12:** PAC-Bayes (average over posterior distribution)



**Unifying insight:** Average-case analysis is tighter than worst-case. Modern theory exploits data structure and algorithmic properties.

✓ **Checkpoint**

**Reflection:** Why is VC dimension still important if Rademacher complexity and PAC-Bayes give tighter bounds?

**Answer:** VC dimension provides a **characterization** of learnability—a hypothesis class is PAC learnable if and only if it has finite VC dimension. Tighter bounds are useful for quantitative predictions, but VC dimension answers the fundamental qualitative question: “is this problem learnable at all?”

### Principle 3: Implicit vs. Explicit Regularization

**Motivation:** One of the most surprising discoveries in modern learning theory is that regularization can emerge from the learning algorithm itself, without explicit penalties in the loss function. This explains why neural networks trained with simple gradient descent—no weight decay, no dropout—still generalize. The algorithmic choices (initialization, learning rate, architecture) encode implicit preferences for certain solutions over others.

Regularization appears in many forms:

← **From Week 5: Ridge Regression**

Ridge regression adds explicit  $\lambda \|\theta\|^2$  penalty to the loss. This was our first systematic approach to controlling complexity: larger  $\lambda$  means simpler models, trading bias for variance.

← **From Week 9: Boosting and Margins**

AdaBoost has **no explicit regularization**, yet it doesn’t overfit even after thousands of iterations. The secret: boosting implicitly maximizes the margin, which acts as a complexity-independent measure of confidence.

← **From Week 10: Follow-The-Regularized-Leader**

FTRL in online learning explicitly adds regularization  $R(\theta)$  to enable logarithmic regret. But even without it, algorithms like Hedge have implicit regularization through the exponential weighting scheme.

| Week | Regularization Type | Mechanism                                                                   |
|------|---------------------|-----------------------------------------------------------------------------|
| 5    | Explicit $\ell_2$   | Ridge regression: $\min \mathcal{L} + \lambda \ \theta\ ^2$                 |
| 6-8  | RKHS norm           | Kernel ridge regression: $\min \mathcal{L} + \lambda \ f\ _{\mathcal{H}}^2$ |
| 9    | Margin maximization | Boosting implicitly maximizes margin                                        |
| 10   | Regularized leader  | FTRL: $\min \mathcal{L} + R(\theta)$                                        |
| 12   | Implicit bias       | GD implicitly minimizes norm                                                |

Table 3: Evolution of regularization across the course

**Unifying insight:** Regularization can be explicit (added to loss) or implicit (emergent from algorithm). Both control complexity and improve generalization.

## Principle 4: The Kernel Connection

**Motivation:** Kernels emerged in Weeks 6-8 as an elegant mathematical framework for nonlinear learning. But their influence extends far beyond those three weeks. Kernels provide a unifying language that connects seemingly disparate algorithms: boosting (coordinate descent in RKHS), neural networks (NTK limit), and even online learning (kernel-based expert algorithms). This universality suggests that kernels capture something fundamental about learning.

Kernels appear throughout the course as a unifying mathematical framework:

### ← From Week 6: Kernel Trick

The kernel trick lets us work in infinite-dimensional feature spaces without ever computing the features explicitly. Computing  $k(x, x') = \langle \varphi(x), \varphi(x') \rangle$  is enough—we never need  $\varphi(x)$  itself!

### ← From Week 7: Representer Theorem

The representer theorem says solutions to regularized ERM in RKHS have the form  $f^* = \sum_i \alpha_i k(x_i, \cdot)$ . This is **not** an algorithm-specific result—it holds for any loss and any RKHS. This universality is powerful.

### ← From Week 9: Boosting as Coordinate Descent

We showed that boosting can be viewed as coordinate descent in an RKHS of decision trees. This connection explains boosting's success and provides a bridge between ensemble methods and kernel methods.

### Connection

#### The Kernel Thread:

- **Week 6:** RKHS theory, reproducing property, kernel trick
- **Week 7:** Kernel ridge regression, representer theorem
- **Week 8:** Kernel PCA, kernel ridge regression, MMD, kernel mean embeddings
- **Week 9:** Boosting as coordinate descent in RKHS
- **Week 12:** Neural Tangent Kernel connects deep learning to kernels

**Unifying insight:** Kernels provide a universal language for learning. They connect linear methods (via feature maps), nonlinear methods (via kernel trick), and even deep learning (via NTK).

## Principle 5: Optimization Matters

**Motivation:** Classical learning theory treats optimization as a solved problem: “assume we find the empirical risk minimizer in  $\mathcal{H}$ .” This assumption is reasonable for convex problems like linear regression or SVMs. But for deep learning, optimization is inseparable from generalization. **What** we optimize (the loss function) matters less than **how** we optimize (the algorithm, initialization, learning rate). This shift from “what” to “how” is one of modern learning theory’s most important insights.

Classical theory treats optimization as solved (assume we find the best hypothesis in  $\mathcal{H}$ ). Modern theory recognizes optimization is crucial:

#### ← From Week 10: Online Learning

Online learning directly analyzes optimization dynamics through regret bounds. There’s no separation between “learning” and “optimization”—they’re the same process. This algorithmic perspective carries over to batch learning.

#### ← From Week 11: Loss Landscape

Week 11 showed that neural network loss landscapes are non-convex with many local minima. Yet SGD reliably finds good solutions. The landscape structure—especially the geometry near initialization—determines what gradient descent learns.

- **Week 10:** Online learning analyzes optimization directly (regret bounds)
- **Week 11:** Gradient descent dynamics, loss landscape, initialization
- **Week 12:** Implicit regularization, edge of stability, grokking

**Unifying insight:** The learning algorithm matters as much as the hypothesis class. Generalization depends on **how** we search, not just **what** we search over.

### Thought Experiment

**Thought experiment:** Consider two algorithms that both minimize empirical risk on the same hypothesis class  $\mathcal{H}$ :

- Algorithm A: Exhaustive search over all  $h \in \mathcal{H}$
- Algorithm B: Gradient descent from random initialization

Classical theory says they should generalize identically (same  $\mathcal{H}$ , same empirical risk minimizer). But Algorithm B generalizes much better! Why?

**Answer:** Algorithm B has implicit regularization. It naturally finds simple solutions near initialization. Algorithm A might find complex, memorizing solutions. The algorithm, not just the hypothesis class, determines generalization.

## Final Reflection

### Author's Take

Statistical learning theory has evolved dramatically over the past three decades:

- **1990s:** PAC learning, VC dimension, worst-case analysis
- **2000s:** Kernel methods, boosting, margin theory
- **2010s:** Deep learning challenges classical theory
- **2020s:** New theory emerges (NTK, benign overfitting, scaling laws)

We're living through a renaissance in learning theory. The gap between theory and practice is narrowing, but many mysteries remain.

The next decade will likely bring:

- Tighter generalization bounds
- Better understanding of transformers and attention
- Unified theory of feature learning
- Predictive theory of emergent abilities

The journey continues. The tools you've learned in this course—PAC learning, kernel methods, optimization theory—will serve as foundation for whatever comes next.

**Thank you for joining this journey through the beautiful mathematics of learning!**

## Practice Problems

### Tier 1 (★): Foundational Understanding

**Problem 1 (★):** Compute the NTK for a two-layer network  $f(x; W, a) = \sum_{j=1}^d a_j \sigma(w_j^T x)$  with fixed output weights  $a_j = \pm 1$  and trainable input weights  $w_j$ . Show that as  $d \rightarrow \infty$ , the NTK converges to a deterministic kernel.

**Problem 2 (★):** Implement double descent on a toy dataset. Generate  $n = 100$  data points from a noisy linear model. Train polynomial regression with degrees  $d = 1, 2, \dots, 150$ . Plot test error vs. degree and identify the interpolation threshold.

**Problem 3 (★):** Verify implicit bias toward max-margin. Train logistic regression on 2D linearly separable data. Track  $\frac{\theta(t)}{\|\theta(t)\|}$  over time and compare to the SVM solution. Does the direction converge?

**Problem 4 (★):** Calculate a PAC-Bayes bound for a simple network. Train a 2-layer network on MNIST. Compute  $\text{KL}(Q \parallel P)$  where  $P$  is the initialization and  $Q$  is a Gaussian around the trained weights. Is the bound non-vacuous?

### Tier 2 (★★): Deeper Analysis

**Problem 5 (★★):** Derive the training dynamics in the NTK regime. Starting from  $d_{\theta}^g t = -\eta \nabla_{\theta} \mathcal{L}$ , show that the network outputs evolve according to  $d_{\theta}^f t = -\eta \Theta(X, X)(f - y)$ . Solve this ODE and show convergence to perfect interpolation.

**Problem 6 (★★):** Study benign overfitting in a toy linear model. Consider  $y = X\beta^* + \varepsilon$  with  $p > n$  and compare the minimum  $\ell_2$ -norm interpolator  $\hat{\beta} = X^+ y$  across different covariance spectra. When does interpolation help or hurt?